

Introduction to R

Session 4 – Data exploration

Statistical Consulting Centre

consulting@stat.auckland.ac.nz
The Department of Statistics
The University of Auckland

19 July, 2017



SCIENCE
DEPARTMENT OF STATISTICS

for loop to get column summary statistics

```
for (i in c("Age.group", "Sex", "Smoke.group", "Race.group",  
  "BMI.group")) {  
  print(i)  
  print(table(combined.df[, i]))  
}
```

```
## [1] "Age.group"
```

```
##
```

```
## Under 35 36 to 60 Over 61
```

```
##      5585      5969      5476
```

```
## [1] "Sex"
```

```
##
```

```
## Female    Male
```

```
##    9077    7953
```

```
## [1] "Smoke.group"
```

```
##
```

```
##      No      Yes
```

for loop to get column summary statistics

```
Categorical.df <- combined.df[, c("Age.group", "Sex",  
  "Smoke.group", "Race.group", "BMI.group")]  
  
for (i in 1:ncol(Categorical.df)) {  
  print(table(Categorical.df[, i]))  
}
```

```
##  
## Under 35 36 to 60 Over 61  
##      5585      5969      5476  
##  
## Female    Male  
##      9077     7953  
##  
##    No    Yes  
##  4255  4371  
##
```

for loop to get column summary statistics

```
Continous.df <- combined.df[, c("Age", "Height", "Weight",  
    "BMI", "Baseline", "PreTrt", "PostTrt")]
```

```
for (i in 1:ncol(Continous.df)) {  
  print(i)  
  print(mean(Continous.df[, i], na.rm = TRUE))  
}
```

```
## [1] 1
```

```
## [1] 48.7919
```

```
## [1] 2
```

```
## [1] 65.43787
```

```
## [1] 3
```

```
## [1] 165.0315
```

```
## [1] 4
```

```
## [1] 27.03084
```

Smart way: (apply) to get column summary statistics

```
apply(X, MARGIN, FUN, ...)
```

- X: A data frame.
- MARGIN: 1 indicates rows, 2 indicates columns.
- FUN: function, what do you want R to do with the rows or columns of the data frame
- ...: optional arguments to FUN.

Translation: Do something (FUN) to every row (or column) (MARGIN) of a data frame (X).

Smart way: (apply) to get column summary statistics

```
apply(Categorical.df, 2, table)
```

```
## $Age.group
```

```
##
```

```
## 36 to 60  Over 61  Under 35
```

```
##      5969      5476      5585
```

```
##
```

```
## $Sex
```

```
##
```

```
## Female    Male
```

```
##    9077    7953
```

```
##
```

```
## $Smoke.group
```

```
##
```

```
##    No    Yes
```

```
## 4255 4371
```

```
##
```

Smart way: (apply) to get column summary statistics

```
apply(Continuous.df, 2, mean, na.rm = TRUE)
```

```
##      Age      Height      Weight      BMI      Baseline
## 48.79190 65.43787 165.03151 27.03084 206.04918
##  PreTrt   PostTrt
## 206.03711 206.03717
```

Smart way: (apply) to get column summary statistics

```
apply(Continuous.df, 2, sd, na.rm = TRUE)
```

```
##      Age      Height      Weight      BMI      Baseline
## 19.720956  3.913986 39.568004  5.833157 44.837071
##   PreTrt   PostTrt
## 45.254414 45.618231
```


`apply()` using self-defined R function

Functions used in `apply()` can be self-defined.

```
na.check <- function(someinput) {  
  test.na <- is.na(someinput)  
  sum(test.na)  
}
```

Take an educated guess at what `na.check()` does.

apply() using a self-defined R function

Let's look at what each row of `na.check()` does.

```
test1 <- Continous.df$BMI[1:10]
test1
```

```
## [1] 25.46094      NA 27.55498 29.36357      NA
## [6] 27.04702 22.56770 20.81527 25.43695 36.90559
```

```
test.na <- is.na(test1)
test.na
```

```
## [1] FALSE TRUE FALSE FALSE TRUE FALSE FALSE
## [8] FALSE FALSE FALSE
```

```
sum(test.na)
```

```
## [1] 2
```

apply() using a self-defined R function

Let's now use `na.check()` in `apply()`.

```
apply(Continuous.df, 2, na.check)
```

```
##      Age      Height      Weight      BMI      Baseline
##      0         40         55         63         968
## PreTrt PostTrt
##      968      968
```

Another way,

```
apply(Continuous.df, 2, function(x) sum(is.na(x)))
```

```
##      Age      Height      Weight      BMI      Baseline
##      0         40         55         63         968
## PreTrt PostTrt
##      968      968
```

A slightly more complicated function

```
mytab <- function(someinput){  
  n <- length(someinput)  
  n.missing <- na.check(someinput)  
  n.complete <- n - n.missing  
  mymean <- round(mean(someinput, na.rm = TRUE), 2)  
  mysd <- round(sd(someinput, na.rm = TRUE), 2)  
  mystder <- round(mysd/sqrt(n.complete), 2)  
  Lower.CI <- round(mymean - 1.96*mystder, 2)  
  Upper.CI <- round(mymean + 1.96*mystder, 2)  
  c(Complete.obs = n.complete, Missing.obs = n.missing,  
     Mean = mymean, Std.Error = mystder,  
     Lower.CI = Lower.CI, Upper.CI = Upper.CI)  
}
```

Take a *more* educated guess at what mytab() does?

A slightly more complicated function

- For the R novice, `mytab()` is possibly terrifying!
- We too were R novices once!
- Our advice on understanding what an R function does?

“Use a data set for input into the function and work through it one line of code at a time.”

- We “experts” still do this!

mytab()

```
apply(Continuous.df, 2, mytab)
```

##	Age	Height	Weight	BMI
## Complete.obs	17030.00	16990.00	16975.00	16967.00
## Missing.obs	0.00	40.00	55.00	63.00
## Mean	48.79	65.44	165.03	27.03
## Std.Error	0.15	0.03	0.30	0.04
## Lower.CI	48.50	65.38	164.44	26.95
## Upper.CI	49.08	65.50	165.62	27.11

##	Baseline	PreTrt	PostTrt
## Complete.obs	16062.00	16062.00	16062.00
## Missing.obs	968.00	968.00	968.00
## Mean	206.05	206.04	206.04
## Std.Error	0.35	0.36	0.36
## Lower.CI	205.36	205.33	205.33
## Upper.CI	206.74	206.75	206.75

More descriptive stats

Calculate the mean Baseline score for male and female patients

```
with(combined.df, mean(Baseline[Sex == "Male"],  
                        na.rm = TRUE))
```

```
## [1] 203.637
```

```
with(combined.df, mean(Baseline[Sex == "Female"],  
                        na.rm = TRUE))
```

```
## [1] 208.1786
```

Better way

```
with(combined.df, tapply(Baseline, Sex, mean, na.rm = T))
```

```
##      Female      Male  
## 208.1786 203.6370
```


tapply()

```
with(issp.df, tapply(total.lik, Income, mean, na.rm = TRUE))
```

```
tapply(X, INDEX, FUN, ...)
```

Translation: Apply function FUN to X for each level in the grouping factor INDEX

How about calculating the mean Baseline score for each Race group (3 levels)?

tapply()

```
with(combined.df, tapply(Baseline, Race.group, mean,  
  na.rm = TRUE))
```

```
##    African Caucasian      Other  
## 201.5053 208.1396 199.9884
```

tapply()

```
with(combined.df, tapply(Baseline, Race.group, mytab))  
## $African  
## Complete.obs Missing.obs      Mean Std.Error  
##      4427.00      433.00    201.51      0.68  
##      Lower.CI      Upper.CI  
##      200.18      202.84  
##  
## $Caucasian  
## Complete.obs Missing.obs      Mean Std.Error  
##     11111.00      501.00    208.14      0.42  
##      Lower.CI      Upper.CI  
##      207.32      208.96  
##  
## $Other  
## Complete.obs Missing.obs      Mean Std.Error  
##      519.00      34.00    199.99      1.90  
##      Lower.CI      Upper.CI  
##      196.27      203.71
```

Data cleaning

```
str(combined.df)
```

```
## 'data.frame':    17030 obs. of  13 variables:
## $ ID             : int  3 4 9 10 11 19 34 44 45 48 ...
## $ Age            : int  21 32 48 35 48 44 42 24 67 56 ...
## $ Sex            : Factor w/ 2 levels "Female","Male": 2 1 1 2 ...
## $ Weight         : num  180 NA 150 204 155 ...
## $ Height         : num  70.4 63.9 61.8 69.8 NA 70.2 62.6 64.4 ...
## $ Smoke.group    : Factor w/ 2 levels "No","Yes": NA NA 1 NA 1 ...
## $ Race.group     : Factor w/ 3 levels "African","Caucasian",...
## $ BMI            : num  25.5 NA 27.6 29.4 NA ...
## $ BMI.group      : Factor w/ 3 levels "normal","overweight",...
## $ Age.group      : Factor w/ 3 levels "Under 35","36 to 60",...
## $ Baseline       : int  268 160 236 225 260 187 216 137 NA 156 ...
## $ PreTrt         : int  276 170 245 231 256 194 212 135 NA 157 ...
## $ PostTrt        : int  281 170 252 235 257 195 222 136 NA 159 ...
```

Data cleaning

```
str(combined.df[, 11:13])
```

```
## 'data.frame':    17030 obs. of  3 variables:
## $ Baseline: int   268 160 236 225 260 187 216 137 NA 156 . .
## $ PreTrt  : int   276 170 245 231 256 194 212 135 NA 157 . .
## $ PostTrt : int   281 170 252 235 257 195 222 136 NA 159 . .
```

- a common problem is the dataset where some of the column names are not names of the variables, but *values* of variable.
- the last three columns are Serum Cholesterol levels, mg/100ml, measured on: Day 1, Day 5 and Day 10, which we have recoded to Baseline, PreTrt and PostTrt, respectively.

tidyr R package

```
library(tidyr)
```

Two main functions: - `gather()` - takes multiple columns and combines based at a *key* value, i.e. Time. - `spread()` - opposite of `gather()`

gather()

```
combined.long.df <- gather(combined.df, "Baseline", "PreTrt",  
  "PostTrt", key = Time, value = Cholesterol)
```

gather()

```
combined.long.df <- gather(combined.df, key = Time, value= Cholesterol,
                           -ID, -Age, -Age.group, -Sex, -Weight,
                           -Smoke.group, -Race.group, -BMI,
                           -BMI.group)
```


gather()

```
head(combined.long.df)
```

##	ID	Age	Sex	Weight	Height	Smoke.group	Race.group
## 1	3	21	Male	179.5	70.4	<NA>	Caucasian
## 2	4	32	Female	NA	63.9	<NA>	Caucasian
## 3	9	48	Female	149.7	61.8	No	Caucasian
## 4	10	35	Male	203.5	69.8	<NA>	Caucasian
## 5	11	48	Male	155.3	NA	No	Caucasian
## 6	19	44	Male	189.6	70.2	Yes	African

##	BMI	BMI.group	Age.group	Time	Cholesterol
## 1	25.46094	overweight	Under 35	Baseline	268
## 2	NA	<NA>	Under 35	Baseline	160
## 3	27.55498	overweight	36 to 60	Baseline	236
## 4	29.36357	overweight	Under 35	Baseline	225
## 5	NA	<NA>	36 to 60	Baseline	260
## 6	27.04702	overweight	36 to 60	Baseline	187

gather()

```
str(combined.long.df)
```

```
## 'data.frame':    51090 obs. of  12 variables:
## $ ID           : int   3  4  9 10 11 19 34 44 45 48 ...
## $ Age          : int   21 32 48 35 48 44 42 24 67 56 ...
## $ Sex          : Factor w/ 2 levels "Female","Male": 2 1 1 2 ...
## $ Weight       : num   180 NA 150 204 155 ...
## $ Height       : num   70.4 63.9 61.8 69.8 NA 70.2 62.6 64.4 ...
## $ Smoke.group  : Factor w/ 2 levels "No","Yes": NA NA 1 NA 1 ...
## $ Race.group   : Factor w/ 3 levels "African","Caucasian",...
## $ BMI          : num   25.5 NA 27.6 29.4 NA ...
## $ BMI.group    : Factor w/ 3 levels "normal","overweight",...
## $ Age.group    : Factor w/ 3 levels "Under 35","36 to 60",...
## $ Time         : chr    "Baseline" "Baseline" "Baseline" "Base
## $ Cholesterol : int   268 160 236 225 260 187 216 137 NA 156
```

spread()

```
combined.wide.df <- spread(combined.long.df, key = Time, value
```

dplyr R package

```
Patient.df <- read.csv("../..\\data\\Patient.csv", stringsAsFactors = FALSE)

step1 <- mutate(Patient.df, Sex = as.factor(Sex), Smoke.group = factor(ifelse(Sex == 1, "Yes", "No")), Race.group = factor(ifelse(Race == 1, "Caucasian", "African", "Other")), BMI = (Weight/Height^2))

step2 <- select(step1, -Smoke, -Race)

step3 <- mutate(step2, BMI.group = factor(ifelse(BMI >= 30, "obese", ifelse(BMI >= 25, "overweight", "normal")), levels = c("normal", "overweight", "obese")))

step4 <- mutate(step3, Age.group = factor(ifelse(Age <= 35, "Under 36", ifelse(Age <= 60, "36 to 60", "Over 61")), levels = c("Under 36", "36 to 60", "Over 61")))
```

Piping %>%

$x \%>\% f(y)$ is equivalent to $f(x, y)$, where x is typically the data-frame.

Piping operator %>%

```
combine.wide.df <- Patient.df %>% mutate(  
  Sex = as.factor(Sex),  
  Smoke.group = factor(ifelse(Smoke == 1, "Yes", "No")),  
  Race.group = factor(ifelse(  
    Race == 1, "Caucasian",  
    ifelse(Race == 2, "African", "Other")  
  )),  
  BMI = (Weight / Height ^ 2) * 703  
) %>%  
select(-Smoke, -Race) %>%  
mutate(BMI.group = factor(  
  ifelse(BMI >= 30, "obese",  
    ifelse(BMI >= 25, "overweight",  
      "normal")),  
  levels = c("normal", "overweight", "obese")  
) %>%
```

Piping operator %>%

```
mutate(Age.group = factor(
  ifelse(Age <= 35, "Under 35",
    ifelse(Age <= 60, "36 to 60", "Over 61")),
  levels = c("Under 35", "36 to 60", "Over 61")
)) %>%
left_join(Cholesterol.df) %>% rename(
  ID = Patient.ID,
  Baseline = Day1,
  PreTrt = Day5,
  PostTrt = Day10
)
```

```
combined.long.df <- combine.wide.df %>%
  gather("Baseline", "PreTrt", "PostTrt",
    key = Time, value = Cholesterol)
```

Grouping

```
with(combined.df, tapply(Baseline, Race.group, mean, na.rm = T
```

```
##      African Caucasian      Other
## 201.5053 208.1396 199.9884
```

```
combine.wide.df %>% group_by(Race.group) %>% summarise(BaseMean =
  na.rm = TRUE))
```

```
## # A tibble: 4 x 2
##   Race.group BaseMean
##   <fctr>      <dbl>
## 1 African 201.5053
## 2 Caucasian 208.1396
## 3 Other 199.9884
## 4 NA 213.0000
```


Grouping

```
with(combined.df, tapply(Baseline, Race.group, mytab))
```

```
## $African
```

## Complete.obs	Missing.obs	Mean	Std.Error
## 4427.00	433.00	201.51	0.68
## Lower.CI	Upper.CI		
## 200.18	202.84		

```
##
```

```
## $Caucasian
```

## Complete.obs	Missing.obs	Mean	Std.Error
## 11111.00	501.00	208.14	0.42
## Lower.CI	Upper.CI		
## 207.32	208.96		

```
##
```

```
## $Other
```

## Complete.obs	Missing.obs	Mean	Std.Error
## 510.00	24.00	100.00	1.00

Grouping

```
combine.wide.df %>% group_by(Age.group, Race.group, BMI.group)
  count()
```

```
## # A tibble: 38 x 4
```

```
## # Groups:   Age.group, Race.group, BMI.group [38]
```

	Age.group	Race.group	BMI.group	n
	<fctr>	<fctr>	<fctr>	<int>
## 1	Under 35	African	normal	906
## 2	Under 35	African	overweight	517
## 3	Under 35	African	obese	470
## 4	Under 35	African	NA	1
## 5	Under 35	Caucasian	normal	1800
## 6	Under 35	Caucasian	overweight	1031
## 7	Under 35	Caucasian	obese	602
## 8	Under 35	Caucasian	NA	10
## 9	Under 35	Other	normal	139
## 10	Under 35	Other	overweight	67

Grouping

```
combine.wide.df %>% group_by(Age.group, Race.group, BMI.group)
  count() %>% arrange(desc(n))
```

```
## # A tibble: 38 x 4
```

```
## # Groups:   Age.group, Race.group, BMI.group [38]
```

##	Age.group	Race.group	BMI.group	n
##	<fctr>	<fctr>	<fctr>	<int>
## 1	Under 35	Caucasian	normal	1800
## 2	Over 61	Caucasian	overweight	1715
## 3	Over 61	Caucasian	normal	1674
## 4	36 to 60	Caucasian	overweight	1422
## 5	36 to 60	Caucasian	normal	1300
## 6	36 to 60	Caucasian	obese	1143
## 7	Under 35	Caucasian	overweight	1031
## 8	Under 35	African	normal	906
## 9	Over 61	Caucasian	obese	882
## 10	36 to 60	African	obese	656

Grouping

```
combine.wide.df %>% group_by(Age.group, Race.group, BMI.group)
count() %>% spread(Age.group, n)
```

```
## # A tibble: 14 x 5
```

```
## # Groups:   Race.group, BMI.group [14]
```

```
##   Race.group BMI.group `Under 35` `36 to 60` `Over 61`
## *   <fctr>      <fctr>      <int>      <int>      <int>
## 1   African    normal        906        580        372
## 2   African overweight      517        654        375
## 3   African    obese        470        656        313
## 4   African      NA           1           3         13
## 5 Caucasian    normal      1800       1300       1674
## 6 Caucasian overweight    1031       1422       1715
## 7 Caucasian    obese        602       1143        882
## 8 Caucasian      NA          10           7         26
## 9    Other     normal       139          88         55
## 10   Other overweight      67          57         25
```

Grouping

```
combine.wide.df %>% group_by(Age.group, Race.group, BMI.group)
  count() %>% spread(Age.group, n) %>% filter(!is.na(BMI.group),
  !is.na(Race.group))
```

```
## # A tibble: 9 x 5
```

```
## # Groups:   Race.group, BMI.group [9]
```

	Race.group	BMI.group	`Under 35`	`36 to 60`	`Over 61`
	<fctr>	<fctr>	<int>	<int>	<int>
## 1	African	normal	906	580	372
## 2	African	overweight	517	654	375
## 3	African	obese	470	656	313
## 4	Caucasian	normal	1800	1300	1674
## 5	Caucasian	overweight	1031	1422	1715
## 6	Caucasian	obese	602	1143	882
## 7	Other	normal	139	88	55
## 8	Other	overweight	67	57	25
## 9	Other	obese	28	57	24

Summary

- `apply()`
- `tapply()`
- `gather()` and `spread()` of `tidyr` R package
- `dplyr` R package
- piping operator `%>%`